

# Deep Learning with R

Chenshu Liu

April 2022

## Contents

<b>1</b>	<b>Dense Layer Model</b>	<b>2</b>
1.1	Data & Library . . . . .	2
1.2	Data Preprocessing . . . . .	2
1.3	Train Test Split . . . . .	2
1.4	Modeling . . . . .	3
<b>2</b>	<b>Regularization</b>	<b>6</b>
2.1	Data & Library . . . . .	6
2.2	Data Preprocessing . . . . .	6
2.3	L2 Regularization Model . . . . .	7
2.4	Dropout Regularization Model . . . . .	8
<b>3</b>	<b>Neural Network Optimizations</b>	<b>11</b>
3.1	Data & Library . . . . .	11
3.2	Data Preprocessing . . . . .	11
3.3	Modeling . . . . .	12
<b>4</b>	<b>Convolution Neural Network</b>	<b>14</b>
4.1	Data & Library . . . . .	14
4.2	Data Preprocessing . . . . .	14
4.3	Modeling . . . . .	14

# 1 Dense Layer Model

## 1.1 Data & Library

```
# import spreadsheet files
library(readr)
```

```
## Warning: replacing previous import 'lifecycle::last_warnings' by
## 'rlang::last_warnings' when loading 'hms'
```

```
## Warning: replacing previous import 'lifecycle::last_warnings' by
## 'rlang::last_warnings' when loading 'tibble'
```

```
## Warning: replacing previous import 'lifecycle::last_warnings' by
## 'rlang::last_warnings' when loading 'pillar'
```

```
# deep learning package
library(keras)
```

```
## Warning: package 'keras' was built under R version 4.1.2
```

```
setwd("~/Documents/Programming/R/Deep Learning with R/Datasets")
data <- read_csv("SimulatedBinaryClassificationDataset.csv",
                 col_names = TRUE)
```

```
## Rows: 50000 Columns: 11
## -- Column specification -----
## Delimiter: ","
## dbl (11): Var1, Var2, Var3, Var4, Var5, Var6, Var7, Var8, Var9, Var10, Target
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

## 1.2 Data Preprocessing

```
# data.frame --> matrix
data <- as.matrix(data)
# remove the row and col names, leaving only numerical values
dimnames(data) = NULL
mode(data)
```

```
## [1] "numeric"
```

## 1.3 Train Test Split

```

# train and test split index
set.seed(123)
index <- sample(2,
               nrow(data),
               replace = TRUE,
               prob = c(0.9, 0.1))

# data splitting
x_train <- data[index == 1, 1:10]
x_test <- data[index == 2, 1:10]
y_test_actual <- data[index == 2, 11]

# use the to_categorical function in keras package for one-hot encoding
y_train <- to_categorical(data[index == 1, 11])

```

```
## Loaded Tensorflow version 2.8.0
```

```
y_test <- to_categorical(data[index == 2, 11])
```

## 1.4 Modeling

```

model <- keras_model_sequential() %>%
  # layer_dense means a densely connected layer
  layer_dense(name = "DeepLayer1",
              units = 10, # hyperparameter: the number of nodes
              activation = "relu",
              # the first layer need to have specification about the input dimension
              input_shape = c(10)) %>%
  layer_dense(name = "DeepLayer2",
              units = 10,
              activation = "relu") %>%
  layer_dense(name = "OutputLayer",
              units = 2,
              # softmax function will provide probabilities of the nodes
              activation = "softmax")

```

```
summary(model)
```

```
## Model: "sequential"
```

```
## -----
```

## Layer (type)	Output Shape	Param #
## =====		
## DeepLayer1 (Dense)	(None, 10)	110
## DeepLayer2 (Dense)	(None, 10)	110
## OutputLayer (Dense)	(None, 2)	22
## =====		

```

## Total params: 242
## Trainable params: 242


```

```
## Non-trainable params: 0
```

```
## -----
```

```
model %>% compile(  
  # another way to calculate loss, besides mean-squared-error  
  loss = "categorical_crossentropy",  
  # a special way of gradient descent  
  optimizer = "adam",  
  # measurement of model performance - using accuracy to measure  
  metrics = c("accuracy"))  
  
history <- model %>%  
  fit(x_train,  
    y_train,  
    # number of full forward & backward propagation  
    # (i.e. run 10 times back and forth of all samples)  
    epoch = 10,  
    # instead of propagating the whole dataset at one go, use smaller batches  
    batch_size = 256,  
    # splitting the training set to test itself during training  
    validation_split = 0.1,  
    verbose = 2)  
  
# plot the training history  
plot(history)
```

```
## 'geom_smooth()' using formula 'y ~ x'
```



```
model %>%
  evaluate(x_test,
           # NOTE: here we are still using the one-hot encoded y_test
           y_test)
```

```
##      loss  accuracy
## 0.1493152 0.9612784
```

```
# form predictions
pred <- model %>%
  predict(x_test) %>%
  k_argmax()
# reference for converting tensor to R data types
# https://torch.mlverse.org/technical/tensors/
pred <- as.array(pred)
table(Predicted = pred,
      # NOTE: for confusion matrix, we are using the original y_test_actual, not encoded
      Actual = y_test_actual)
```

```
##      Actual
## Predicted    0    1
##           0 2419 150
##           1   39 2273
```

## 2 Regularization

### 2.1 Data & Library

```
# Applying regularization to deal with overfitting
library(keras)
library(readr)
library(tidyr)
library(tibble)
library(plotly)

## Loading required package: ggplot2

##
## Attaching package: 'plotly'

## The following object is masked from 'package:ggplot2':
##
##   last_plot

## The following object is masked from 'package:stats':
##
##   filter

## The following object is masked from 'package:graphics':
##
##   layout

# specify the number of feature variables for the dataset to be downloaded
num_words <- 5000
imdb <- dataset_imdb(num_words = num_words)

# train test split
c(train_data, train_labels) %<-% imdb$train
c(test_data, test_labels) %<-% imdb$test
```

### 2.2 Data Preprocessing

```
# multi-hot encoding
multi_hot_sequences <- function(sequences, dimension){
  multi_hot <- matrix(0,
    # the number of samples in the sequences
    # sequences are stored as lists
    nrow = length(sequences),
    ncol = dimension)
  for(i in 1 : length(sequences)){
    # sequences[[i]] extracts the label of the words in the text sample i
    # which ever word is included in that sequence will be assigned 1 at row i
    multi_hot[i, sequences[[i]]] <- 1
  }
}
```

```

    }
    multi_hot
  }

train_data <- multi_hot_sequences(train_data, num_words)
test_data <- multi_hot_sequences(test_data, num_words)

```

## 2.3 L2 Regularization Model

```

l2_model <-
  keras_model_sequential() %>%
  layer_dense(units = 16, activation = "relu", input_shape = num_words,
    # apply regularization in the layer_dense function's argument
    kernel_regularizer = regularizer_l2(l = 0.001)) %>%
  layer_dense(units = 16, activation = "relu",
    kernel_regularizer = regularizer_l2(l = 0.001)) %>%
  layer_dense(units = 1, activation = "sigmoid")

l2_model %>% compile(
  optimizer = "adam",
  loss = "binary_crossentropy",
  metrics = list("accuracy")
)

l2_model %>% summary()

```

```

## Model: "sequential_1"
## -----
## Layer (type)                Output Shape          Param #
## =====
## dense_2 (Dense)             (None, 16)            80016
##
## dense_1 (Dense)             (None, 16)            272
##
## dense (Dense)               (None, 1)              17
##
## =====
## Total params: 80,305
## Trainable params: 80,305
## Non-trainable params: 0
## -----

```


```

l2_history <- l2_model %>% fit(
  train_data,
  train_labels,
  epoch = 20,
  batch_size = 512,
  validation_data = list(test_data, test_labels),
  verbose = 2
)

plot(l2_history)

```

```
## 'geom_smooth()' using formula 'y ~ x'
```



## 2.4 Dropout Regularization Model

```
drop_model <- keras_model_sequential() %>%
  layer_dense(units = 16, activation = "relu", input_shape = num_words) %>%
  # a new layer to specify the dropout rate
  layer_dropout(0.6) %>%
  layer_dense(units = 16, activation = "relu") %>%
  layer_dropout(0.6) %>%
  layer_dense(units = 1, activation = "sigmoid")

drop_model %>% compile(
  optimizer = "adam",
  loss = "binary_crossentropy",
  metrics = list("accuracy")
)

drop_model %>% summary()
```

```
## Model: "sequential_2"
```

```
## -----
## Layer (type)                               Output Shape          Param #
## =====
```




```
## dense_5 (Dense)                (None, 16)                80016
##
## dropout_1 (Dropout)            (None, 16)                0
##
## dense_4 (Dense)                (None, 16)                272
##
## dropout (Dropout)              (None, 16)                0
##
## dense_3 (Dense)                (None, 1)                 17
##
## =====
## Total params: 80,305
## Trainable params: 80,305
## Non-trainable params: 0
## -----
```

```
drop_history <- drop_model %>% fit(
  train_data,
  train_labels,
  epoch = 20,
  batch_size = 512,
  validation_data = list(test_data, test_labels),
  verbose = 2
)

plot(drop_history)
```

```
## 'geom_smooth()' using formula 'y ~ x'
```



## 3 Neural Network Optimizations

### 3.1 Data & Library

```
library(readr)
library(keras)

setwd("~/Documents/Programming/R/Deep Learning with R/Datasets")
data.set <- read_csv("RegressionData.csv",
                     col_names = FALSE)

## Rows: 4898 Columns: 11
## -- Column specification -----
## Delimiter: ","
## db1 (11): X1, X2, X3, X4, X5, X6, X7, X8, X9, X10, X11
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

### 3.2 Data Preprocessing

#### 3.2.1 General Preprocessing

```
# transform dataframe to matrix
data.set <- as.matrix(data.set)
# remove column names
dimnames(data.set) <- NULL

# train test split
set.seed(123)
index <- sample(2,
               nrow(data.set),
               replace = TRUE,
               prob = c(0.8, 0.2))

x_train <- data.set[index == 1, 1:10]
x_test <- data.set[index == 2, 1:10]
y_train <- data.set[index == 1, 11]
y_test <- data.set[index == 2, 11]
```

#### 3.2.2 Normalization

```
# normalizing data
mean.train <- apply(x_train, 2, mean)
sd.train <- apply(x_train, 2, sd)
x_train <- scale(x_train)
# use the normalizing parameters from training set to normalize testing set
x_test <- scale(x_test,
               center = mean.train,
               scale = sd.train)
```

### 3.3 Modeling

```
# creating the model
model <- keras_model_sequential() %>%
  layer_dense(units = 25,
              activation = "relu",
              input_shape = c(10)) %>%
  layer_dropout(0.2) %>%
  layer_dense(units = 25,
              activation = "relu") %>%
  layer_dropout(0.2) %>%
  layer_dense(units = 25,
              activation = "relu") %>%
  layer_dropout(0.2) %>%
  layer_dense(units = 1)

model %>% summary()
```

```
## Model: "sequential_3"
## -----
## Layer (type)                Output Shape          Param #
## =====
## dense_9 (Dense)             (None, 25)            275
##
## dropout_4 (Dropout)         (None, 25)            0
##
## dense_8 (Dense)             (None, 25)            650
##
## dropout_3 (Dropout)         (None, 25)            0
##
## dense_7 (Dense)             (None, 25)            650
##
## dropout_2 (Dropout)         (None, 25)            0
##
## dense_6 (Dense)             (None, 1)             26
##
## =====
## Total params: 1,601
## Trainable params: 1,601
## Non-trainable params: 0
## -----
```

```
# compile the model
model %>% compile(
  # the metric for propagation
  loss = "mse",
  optimizer = optimizer_rmsprop(),
  # not for propagation, but for user feedback
  # letting us know the model's performance
  metrics = c("mean_absolute_error"))

# fitting the data
model_history <- model %>%
```


```

fit(x_train,
    y_train,
    epoch = 50,
    batch_size = 32,
    validation_split = 0.1,
    callbacks = c(callback_early_stopping(monitor = "val_mean_absolute_error",
                                          patience = 5)),
    verbose = 2)

plot(model_history)

```

```
## 'geom_smooth()' using formula 'y ~ x'
```



```

# testing the model
c(loss, mae) %<-% (model %>% evaluate(x_test, y_test, verbose = 0))
paste0("Mean Absolute Error on test set is:", mae)

```

```
## [1] "Mean Absolute Error on test set is:0.644970536231995"
```

## 4 Convolution Neural Network

### 4.1 Data & Library

```
library(keras)

# a dataset of numerous images of hand-written numbers from 0-9
mnist <- dataset_mnist()
x_train <- mnist$train$x
y_train <- mnist$train$y
x_test <- mnist$test$x
y_test <- mnist$test$y
```

### 4.2 Data Preprocessing

```
img_row <- dim(x_train)[2]
img_col <- dim(x_test)[3]

# images should have three channels, but the data only has two, need transformation
# because this is gray-scale image, the third channel only has one layer
x_train <- array_reshape(x_train,
                        c(nrow(x_train),
                          img_row,
                          img_col, 1))
x_test <- array_reshape(x_test,
                      c(nrow(x_test),
                        img_row,
                        img_col, 1))
input_shape <- c(img_row, img_col, 1)

# normalize the datasets by dividing the number 255
# because the color gradient is from 0(black) to 255(white)
# dividing by 255 can transform the entries to values between 0 and 1
x_train <- x_train/255
x_test <- x_test/255

# use one-hot encoding to encode the y values
# y values are labels for number 0 to 9, thus we need 10 categories
y_train <- to_categorical(y_train, num_classes = 10)
y_test <- to_categorical(y_test, num_classes = 10)
```

### 4.3 Modeling

```
model <- keras_model_sequential() %>%
  layer_conv_2d(
    # number of filters for transformation
    filters = 16,
    # size of the filters
    kernel_size = c(3,3),
    activation = 'relu',
```

```

    input_shape = input_shape) %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_dropout(rate = 0.25) %>%
  layer_flatten() %>%
  layer_dense(units = 10,
              activation = 'relu') %>%
  layer_dropout(rate = 0.5) %>%
  layer_dense(units = 10,
              # for categorical prediction
              activation = 'softmax')

model %>% summary()

```

```

## Model: "sequential_4"
## -----
## Layer (type)                Output Shape          Param #
## =====
## conv2d (Conv2D)             (None, 26, 26, 16)    160
##
## max_pooling2d (MaxPooling2D) (None, 13, 13, 16)    0
##
## dropout_6 (Dropout)         (None, 13, 13, 16)    0
##
## flatten (Flatten)           (None, 2704)          0
##
## dense_11 (Dense)            (None, 10)            27050
##
## dropout_5 (Dropout)         (None, 10)            0
##
## dense_10 (Dense)            (None, 10)            110
##
## =====
## Total params: 27,320
## Trainable params: 27,320
## Non-trainable params: 0
## -----

```

```

model %>% compile(
  loss = loss_categorical_crossentropy,
  optimizer = optimizer_adadelta(),
  metrics = c('accuracy')
)

model %>% fit(
  x_train,
  y_train,
  batch_size = 128,
  epochs = 12,
  validation_split = 0.2
)

score <- model %>% evaluate(x_test,
                           y_test)

```

score

```
##      loss  accuracy
## 0.1689227 0.9651000
```